

COMMUNITY

BLOGS

PARTNERS

CONTACT US



SUPPORT LOGIN



ENGLISH ▼



English ▼



Products

Solutions

Customers

Services & Support

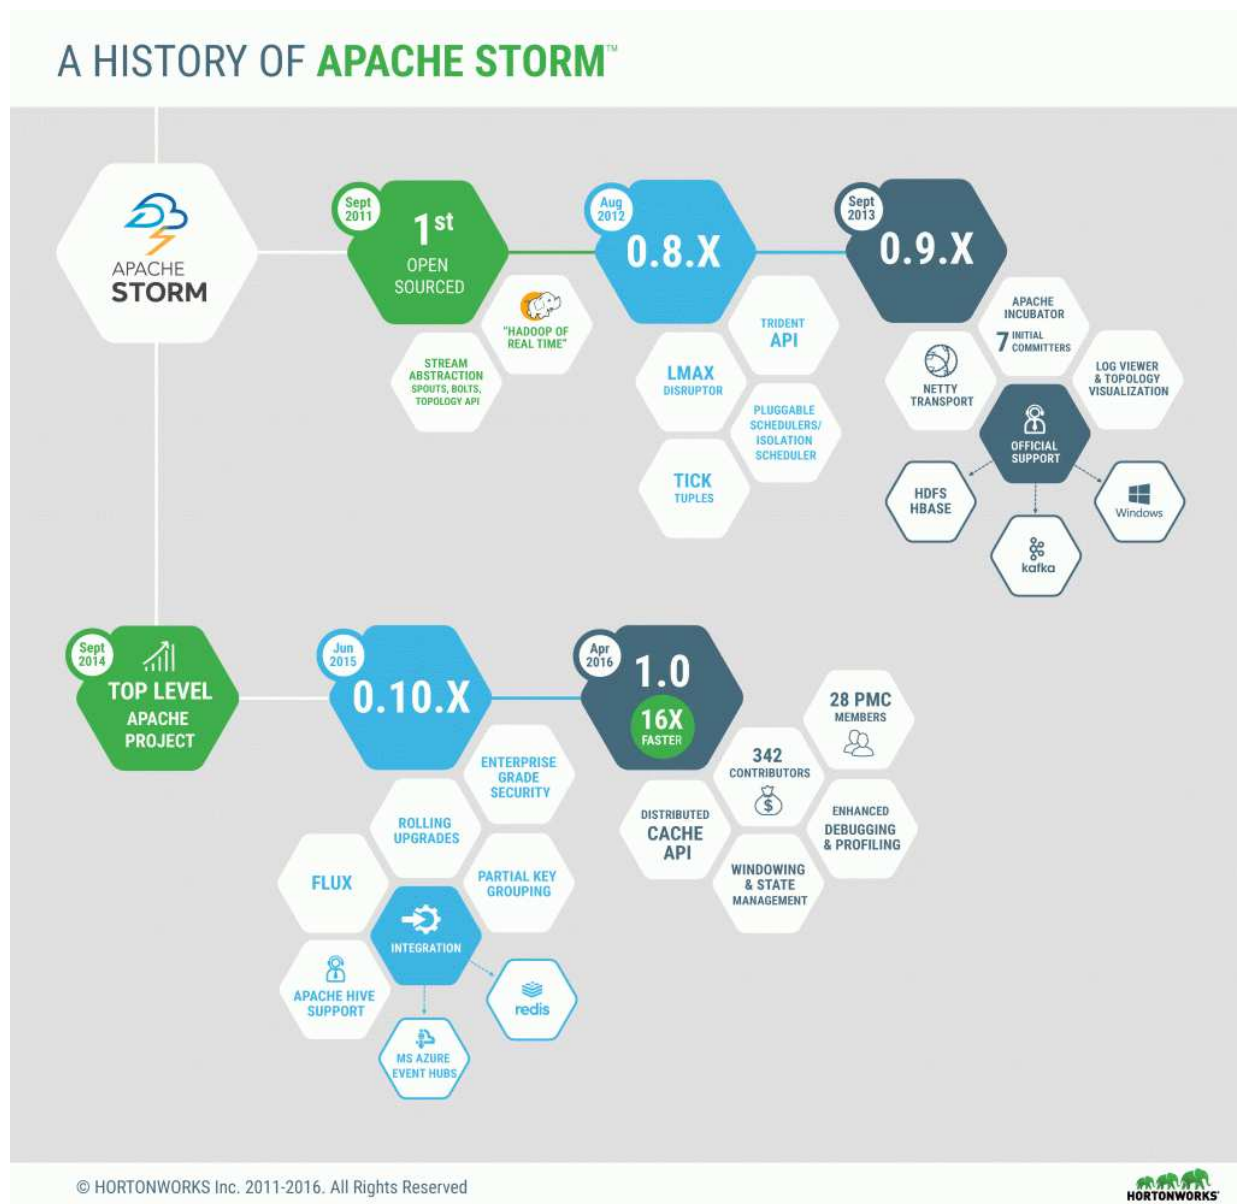
◀ May 18, 2016



A BRIEF HISTORY OF APACHE STORM

by [Guest Author](#)





PART 1: A LITTLE HISTORY

In this series of blog posts, we will provide an in-depth look select features introduced with the release of Apache Storm (Storm) 1.0. To kick off the series, we'll take a look how **Storm** has evolved over the years from its beginnings as an open source project, up to the 1.0 milestone release.

“THE HADOOP OF REAL-TIME”

Storm was originally created by Nathan Marz while he was at Backtype (later acquired by **Twitter**) working on analytics products based on historical and real-time analysis of the Twitter firehose. Nathan envisioned Storm as a

MENU [news-blogs](#)

SHARE



SIGN-UP

Contact Sales?

concept of the “stream” as a distributed abstraction for data in motion, as well as a fault tolerance and reliability model that was difficult, if not impossible, to achieve with a traditional queues and workers architecture.

Nathan open sourced Storm to [GitHub](#) on September 19th, 2011 during his talk at Strange Loop, and it quickly became the most watched JVM project on [GitHub](#). Production deployments soon followed, and the Storm development community rapidly expanded.

At the time [Storm](#) was introduced, Big Data analytics largely involved batch processing in map-reduce on [Apache Hadoop](#) or one of the higher level abstractions like [Apache Pig](#) and Cascading. The introduction of Storm helped drive a change in the way people thought about large scale analytics, spurring the rise of stream processing and real-time analytics.

Early versions of [Storm](#) introduced the familiar stream abstraction and the corresponding Spout/Bolt/Topology API that allowed developers to easily reason about streaming computations. Another feature Storm introduced, and one that to this day remains unique to Storm, is the concept of Distributed Remote Procedure Calls where the inherent parallelism and scalability of Storm can be leveraged in a synchronous, request-response paradigm.

The Storm 0.8.x line of releases introduced the Trident API which added support for exactly-once semantics, micro-batching for increased throughput, stateful processing, and a high-level API for joins, aggregations, grouping, functions, and filters. Other improvements at the time included pluggable schedulers, introduction of the LMAX Disruptor for higher throughput, tick tuples, and improvements to the Storm UI.

STORM MOVES TO APACHE

With encouragement from Andy Feng at Yahoo!, Nathan decided to propose moving Storm to [Apache](#), and the project officially entered the Apache Incubator on September 18, 2013. This move marked the beginning of a fundamental shift in the Storm community away from a model where a single

[MENU](#)[news-blogs](#)[SIGN-UP](#)[Contact Sales?](#)

SHARE



model. The move to [Apache](#) ensured that the project would be governed by a sustainable community, and that no single individual would represent a bottleneck in terms of decision making.

During Storm's time in the Apache Incubator, the 0.9.x line of releases was introduced. In 0.9.x Storm's underlying transport layer based on Omq was replaced with an implementation based on Netty. Not only was the new Netty-based transport almost twice as fast as the previous implementation, but the fact that Netty is a pure Java framework freed users from the requirement of difficult to install, platform specific binaries. Finally, the Netty transport set the stage for authorization and authentication between worker processes.

The 0.9.x line of releases also introduced [Apache Kafka](#) integration as a first class component in the Apache Storm distribution. Prior to this point [Kafka](#) integration was a separate project that had been forked many times, and it was difficult for users to understand which versions were compatible with various versions of Storm and Kafka. Bringing [Kafka](#) integration into the Apache distribution ensured that compatibility was maintained with each [Storm](#) release. Similarly, the [Storm](#) community added support for [HDFS](#) and [Apache HBase](#) integration as well.

Other notable improvements from this time included performance improvements to the Netty transport, pluggable serialization for multi-lang components, topology visualization, a logviewer, and support for Microsoft Windows deployments.

[Apache Storm](#) became a top-level Apache project on September 17, 2014. In the time since first entering the Apache Incubator with a group of 7 initial committers, the [Apache Storm](#) PMC has grown to 28 members strong with contributions from 342 individuals.

ENTERPRISE READINESS

Before graduating from the Apache Incubator, the [Storm community](#) was already working on the next major iteration of the platform: Version 0.10.

[MENU](#)[news-blogs](#)[SIGN-UP](#)[Contact Sales?](#)[SHARE](#)

early days of Apache Hadoop, [Apache Storm](#) originally evolved in an environment where security was not a high-priority concern. Rather, it was assumed that [Storm](#) would be deployed to environments suitably cordoned off from security threats. While a large number of users were comfortable setting up their own security measures for Storm, this proved a hindrance to broader adoption among larger enterprises where security policies prohibited deployment without specific safeguards.

The implementation of enterprise-grade security in [Apache Storm](#) was a momentous effort that involved active collaboration between Yahoo!, Hortonworks, Symantec, and the [broader Apache Storm community](#). Some of the highlights of Storm's security features include:

Kerberos Authentication with Automatic Credential Push and Renewal

Pluggable Authorization and ACLs

Multi-Tenant Scheduling with per-user isolation and configurable resource limits.

User Impersonation

SSL Support for Storm UI, Log Viewer, and DRPC (Distributed Remote Procedure Call)

Secure integration with other Hadoop Projects (such as [ZooKeeper](#), HDFS, [HBase](#), etc.)

User isolation (Storm topologies run as the user who submitted them)

Another important feature of Storm 0.10 was the introduction of Flux, a framework and set of utilities that make defining and deploying Storm topologies less developer-intensive. Prior to Flux, a common complaint from users was that the definition of a topology DAG was often tied up in java code, and any change required recompilation and packaging of the topology. Flux addressed that problem by providing a YAML DSL for defining topologies in a simple text file, decoupling the DAG definition from java code.

MENU

[news-blogs](#)

SIGN-UP



Contact Sales?



SHARE



Some of the key features of Flux include:

Easily configure and deploy Storm topologies (Both Storm core and Micro-batch API) without embedding configuration in your topology code

Support for existing topology code

Define Storm Core API (Spouts/Bolts) using a flexible YAML DSL

YAML DSL support for most Storm components (storm-kafka, storm-hdfs, storm-hbase, etc.)

Convenient support for multi-lang components

External property substitution/filtering for easily switching between configurations/environments (similar to Maven-style `\${variable.name}` substitution)

Development of the 0.10 line of releases also saw a proliferation of additional integration components including JDBC/RDBMS integration, streaming ingest to [Apache Hive](#), [Microsoft Azure Event Hubs](#) integration, and [Redis](#) support. Other important features included rolling upgrade support, logging improvements, and a partial key groupings implementation.

1.0 MILESTONE

On April 12, 2016 the [Apache Storm community](#) announced the release of [Apache Storm 1.0](#), representing yet another major milestone in the evolution of the project. [Version 1.0](#) includes a tremendous number of new features, and usability, management, and performance improvements.

In the coming weeks we will continue this blog series with more in-depth articles covering the important new features included in 1.0, including:

Performance improvements


Windowing and State Management

MENU [news-blogs](#)

SHARE



SIGN-UP 

Contact Sales? 

Management and Debugging improvements

Distributed Cache API

Automatic Backpressure Support

Resource Aware Scheduling


New Integration Components

For a sneak preview of these features, please see the [Apache Storm 1.0 release announcement](#) and stay tuned for more in-depth detail.

Tags:


[HORTONWORKS](#)[HADOOP](#)[HDF](#)[APACHE STORM](#)

COMMENTS

-  **streamanalytix** says:
[March 1, 2017 at 5:44 am](#)
I am extremely impressed with the view point. Thanks for sharing such useful information.

[Apache Spark Streaming](#)

[Reply](#)

-  **sarika** says:
[April 14, 2017 at 3:35 am](#)
Its an awesome post and is quite meaningful.

[Reply](#)

LEAVE A REPLY

Your email address will not be published. Required fields are marked *

Comment

[MENU](#)[news-blogs](#)[SHARE](#)[SIGN-UP](#)[Contact Sales?](#)

Name *

Email *

Website

POST COMMENT

RELATED POSTS

BLOG

9.11.17

Strata Data Conference New York – Powering...

The annual Strata Data Conference will make its next stop in New York the week of September 25-29. The core theme of the conference this year is around driving business transformations...

BLOG

9.7.17

Top 5 Use Case

With the explosion businesses need to strategies and adopt architecture. Build

[VIEW ALL EVENTS](#)

[VIEW ALL WEBCASTS](#)

[Ecosystem](#)

[Press Room](#)

[Videos](#)

[Events](#)

[Webinars](#)

[Be First Insights](#)

[Careers](#)

1.408.675.0983

1.855.8.HORTON

44 (0) 20 3826 1405

Apache, Hadoop, Falcon, Atlas, Tez, Sqoop, Flume, Kafka, Pig, Hive, HBase, Accumulo, Storm, Solr, Spark, Ranger, Knox, Ambari, ZooKeeper, Oozie, Phoenix, NiFi, HAWQ, Zeppelin, Atlas, Slider, Mahout, MapReduce, HDFS, YARN, Metron and the Hadoop elephant and Apache project logos are either registered trademarks or trademarks of the [Apache Software Foundation](#) in the United States or other countries.



© 2011-2017 Hortonworks Inc. All Rights Reserved.

[Privacy Policy](#)

[Terms of Service](#)

[LCA Postings](#)

MENU

[news-blogs](#)

SHARE



SIGN-UP



Contact Sales?

